

Sarcasm or Praise? The Art of Sentiment Analysis with Deep Learning in Film Reviews



Hardefa Rogonondo

Follow

9 min read · Jan 6, 2024



Photo by [riis riiiis](#) on [Unsplash](#)

[GitHub](#)

Chapter 1: Decoding the Audience's Applause (or Boos)

Every time a movie ends, the audience leaves behind a piece of their mind in the reviews they write — applauding stellar performances, critiquing plot holes, or sometimes just ranting. But within this sea of opinions lies valuable insights, not just for movie buffs but for film studios, critics, and streaming platforms that live and breathe audience sentiment.

Here's where we roll out the red carpet for sentiment analysis. It's a nifty AI-powered tool that dives into the text of reviews and surfaces with a clear "yay" or "nay." This isn't just a black-and-white classification gig — it's about deciphering the shades of gray in between. Is it a rave review with a tiny gripe about the CGI, or a scathing critique with a nod to an actor's performance? Our engine's got to figure it out.

But why is this project binge-worthy? Because understanding sentiment is understanding people. It's what drives choices — from which movie gets the weekend slot to what script gets greenlit. And in today's world where data is the new box office hit, this project stands out as a director's cut of the tech world.

So, how did we script this blockbuster? We called in a specialist: the LSTM model. Think of it as the Daniel Day-Lewis of algorithms — versatile, powerful, and capable of getting to the heart of complex narratives. With a dataset sourced from IMDb, as varied as the genres on Netflix, our model learned to read between the lines, picking up on cues as subtle as a nuanced performance in a character-driven indie flick.

Chapter 2: Previous Episodes: The Story So Far

As we dive into the complex world of sentiment analysis with deep learning, it's essential to acknowledge the milestones and the scholarly works that have paved the way for innovations like ours. This journey from rudimentary text analysis to sophisticated AI interpretation is nothing short of a narrative arc in a compelling drama.

Sentiment analysis began as a relatively straightforward task. Early studies focused on identifying positive and negative words within a text. A paper that stands out from this era is Pang and Lee's "Thumbs up? Sentiment Classification using Machine Learning Techniques" (2002), which explored various methods for classifying the sentiment of movie reviews.

The advent of deep learning marked a paradigm shift. Suddenly, computers could understand text in a way that was nuanced and contextually rich. One of the seminal papers, "Distributed Representations of Words and Phrases and their Compositionality" by Mikolov et al. (2013), introduced word embeddings, a breakthrough in how machines could understand text. Another cornerstone was "Sequence to Sequence Learning with Neural Networks" by Sutskever et al. (2014), which revolutionized how models dealt with sequences, a fundamental aspect of sentiment analysis.

Long Short-Term Memory networks (LSTMs) became the stars of the show due to their ability to remember information for long periods. Hochreiter & Schmidhuber's "Long Short-Term Memory" (1997) paper laid the groundwork, while later adaptations demonstrated their efficacy in sentiment analysis. For instance, "Bidirectional LSTM-CRF Models for Sequence Tagging" by Huang et al. (2015) showed the power of LSTMs in understanding context in text.

More recently, the introduction of models like BERT (Bidirectional Encoder Representations from Transformers) in the paper “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding” by Devlin et al. (2018) has pushed the boundaries even further. These models can understand the sentiment of a piece of text with an almost human-like subtlety.

Specific to our realm, several studies have focused on movie review datasets. For instance, “Learning Word Vectors for Sentiment Analysis” by Maas et al. (2011) specifically deals with IMDb reviews, showcasing the potential of machine learning in this domain.

As we stand on the shoulders of these giants, our project doesn't just replicate but innovates, weaving together the threads of previous research into a tapestry that represents the current state-of-the-art in sentiment analysis.

Chapter 3: Crafting the Lens: Our Toolkit and Spells

In any great production, the tools and techniques used behind the scenes are as important as the actors on the stage. In this chapter, we delve into the nuts and bolts of our IMDb Sentiment Analysis Engine, focusing on the rich dataset that forms our foundation, the sophisticated LSTM model that's the brain of our operation, and the hyperparameter tuning that fine-tunes our performance.

Our journey begins with the IMDb dataset, a treasure trove of movie reviews that's as diverse as the cinematic world itself. Each review is a window into a viewer's thoughts and feelings, expressed in natural, unstructured language. This dataset isn't just large; it's a complex, nuanced collection of human


expression, making it both a challenging and an ideal candidate for deep learning-based sentiment analysis.

The star of our show is the LSTM (Long Short-Term Memory) network. Born from the need to understand sequences and remember context over long periods, LSTMs are uniquely suited to text. They're not just reading words; they're grasping the storyline of the sentences, understanding the build-up of emotions. This ability to remember and learn from the context makes LSTMs adept at distinguishing genuine sentiment from mere words.

But even the best actors need direction to give their best performance. That's where hyperparameter tuning comes in, and we've chosen Hyperband as our director. Hyperband is a cutting-edge optimization algorithm that quickly finds the best 'settings' for our model, much like a director finds the best angle for a shot. It's a resource-efficient way to test various configurations and ensure our LSTM is not just good, but great.

Before any text can take the stage, it undergoes a transformation process. We clean it, break it down into understandable pieces, and convert it into a format that our LSTM can understand. This preprocessing step is like rehearsals before the main act, ensuring every word is in its place, ready for the performance.

Deep Learning Framework: The code uses Keras, a high-level neural networks API capable of running on top of TensorFlow, CNTK, or Theano. Keras is widely used for deep learning applications.



Model Architecture: The Sequential model is being used, which is a linear stack of layers common in many deep learning models.

Get Hardefa Rogonondo's stories in your inbox

Join Medium for free to get updates from this writer.

Enter your email


Subscribe

☐ Remember me for faster sign in

Layers:

- **Embedding Layer:** This is often used in deep learning models dealing with text data. It creates word vectors of the text data and is a common representation learning technique.

LSTM Recurrent Unit



- **LSTM (Long Short-Term Memory) Layer:** This is a special kind of RNN (Recurrent Neural Network) capable of learning long-term dependencies. LSTM is a classic deep learning structure particularly suited for sequential data like text.
 - **Dense Layer:** This is a standard neural network layer where each input node is connected to each output node.
1. **Loss Function & Optimizer:** The model uses binary crossentropy as a loss function and the Adam optimizer, both of which are common in deep learning for binary classification tasks.
 2. **Early Stopping:** This is a form of regularization used to prevent overfitting, a common issue in deep learning models.

In this chapter, we've seen the tools and techniques that form the backbone of our sentiment analysis engine. From the raw, unstructured text of IMDb reviews to the finely tuned LSTM model, every component plays a crucial role. As we move forward, these elements come together in a symphony of technology and data, driving us towards our goal of understanding the sentiment behind every review.

Chapter 4: Showtime: The Engine's Big Reveal

After meticulous preparation, it's time to raise the curtain and unveil the centerpiece of our project. This chapter is about the culmination of our efforts — the final, polished version of the IMDb Sentiment Analysis Engine. We'll explore the intricacies of our chosen model, the seamless interplay between the front end and back end, and the thoughtful containerization that makes our system robust and versatile.

Our LSTM model has taken its final form, a well-tuned machine, ready to interpret the emotions of moviegoers. After numerous rehearsals (training sessions) and directorial tweaks (hyperparameter tuning), it stands ready to perform. It's not just any model; it's one that understands the subtleties of human language, can catch the underlying sentiment, and even pick up on the tricky nuances like sarcasm or mixed feelings.



No show can go on without the stage and the crew. Our front end is the stage, a user-friendly interface where users can interact with the model, inputting reviews and receiving sentiment analysis in return. It's designed for simplicity and clarity, ensuring every user, regardless of technical background, can easily understand and use our tool. The back end is our diligent crew, working tirelessly behind the scenes. It's the server, the database, and all the internal logic that processes the data, runs the model, and delivers the results.

In the world of software, being adaptable and portable is crucial. That's where containerization comes in. By packaging our application and its environment into containers, we've made our sentiment analysis engine compact, self-sufficient, and ready to be deployed anywhere. Whether it's a different machine, a different platform, or even a different part of the world, our system is ready to go on tour, bringing sentiment analysis to wherever it's needed.

When the lights dim and the model takes center stage, it's a moment of truth. We see it in action, processing reviews, and delivering insights. The

performance metrics tell us how well it's doing, but the real test is in the user experience. Does it provide valuable, accurate sentiment analysis? Does it help users understand their audience better? This is the ultimate test of our system's success.

In this chapter, we've peeled back the curtain to reveal the final product of our hard work and innovation. Our IMDb Sentiment Analysis Engine, powered by a sophisticated LSTM model and housed within a robust, user-friendly system, is not just a tool; it's a testament to the power of deep learning in understanding human sentiment. As we move forward, we'll continue to refine, improve, and adapt, but for now, we take a moment to appreciate the show we've put on.

Chapter 5: The Closing Credits: Conclusions and Future Directions

As our journey through the realms of sentiment analysis and deep learning draws to a close, it's time to reflect on the path we've traversed and gaze ahead at the roads yet to be explored. In this final chapter, we'll summarize the achievements, lessons learned, and envision the future enhancements that can propel our IMDb Sentiment Analysis Engine to new heights.

- **Curtain Call — Achievements:** Our project set out to capture the pulse of moviegoers through their reviews, and we did just that. By harnessing the power of LSTM networks and deep learning, we've developed an engine that doesn't just read text; it understands sentiment. It's a system that's both powerful in its analytical capabilities and user-friendly, making advanced sentiment analysis accessible to all.
- **Intermission Thoughts — Lessons Learned:** Every production has its challenges, and ours was no different. We learned the importance of

meticulous data preprocessing, the nuances of model tuning, and the criticality of a user-centric design. Perhaps the most significant lesson was that in the realm of AI, the learning never stops — both for the models and the developers.

- **Next Season's Teaser — Future Works:** As proud as we are of our current system, the world of AI and sentiment analysis is ever-evolving, and so are our aspirations. Here's what we're dreaming up next:
- **Expanding the Script:** Incorporating more languages and dialects to make our engine universally applicable.
- **Emotional Depth:** Moving beyond positive and negative to detect a spectrum of emotions for a more nuanced analysis.
- **Real-Time Reviews:** Developing capabilities for real-time analysis, providing instant feedback on newly released movies or trailers.
- **User Feedback Loop:** Implementing a mechanism for users to provide feedback on predictions, continually refining our model's accuracy.
- **Cross-Industry Adaptation:** Customizing the model for different industries, from product reviews to social media monitoring.

As the credits roll on our IMDb Sentiment Analysis Engine, we're reminded of the power of language and technology combined. We've built more than just a tool; we've created a gateway to understanding human emotions at scale. The future is bright and brimming with potential, and we're just getting started. Here's to the next chapter in deep learning and sentiment analysis, where every word and every emotion counts.

References

[1] [Sentiment Analysis Based on Deep Learning: A Comparative Study](#)

[2] Systematic reviews in sentiment analysis: a tertiary study.



Written by Hardefa Rogonondo

13 followers · 5 following

Follow

No responses yet



Write a response

What are your thoughts?

